

Lecture 02

Digital Image Basics

2024-02-01

Sébastien Valade



UNIVERSIDAD NACIONAL
AUTÓNOMA DE
MÉXICO

1. What is a digital image?

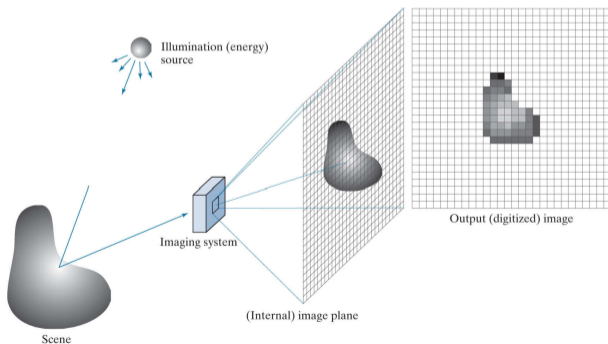
1. image acquisition
2. sampling and quantization
3. 3D projection on 2D plane
4. color image
5. color spaces
6. image histogram

2. Point operations

3. Image processing levels


4. Image manipulation with Python

1. energy from an **illumination source** is reflected from a **scene**
2. the **imaging system** collects the incoming energy and focuses it onto an **image plane**
NB: light-sensing instruments typically use 2-D arrays of photosensors to record incoming light intensity $I(x)$: the CCD (*Charge-Coupled Device*)
3. the image plane is sampled and quantized to produce a **digital image**




1.1. image acquisition

1. energy from an **illumination source** is reflected from a **scene**
2. the **imaging system** collects the incoming energy and focuses it onto an **image plane**
*NB: light-sensing instruments typically use 2-D arrays of photosensors to record incoming light intensity $I(x)$: the CCD (*Charge-Coupled Device*)*
3. the image plane is sampled and quantized to produce a **digital image**



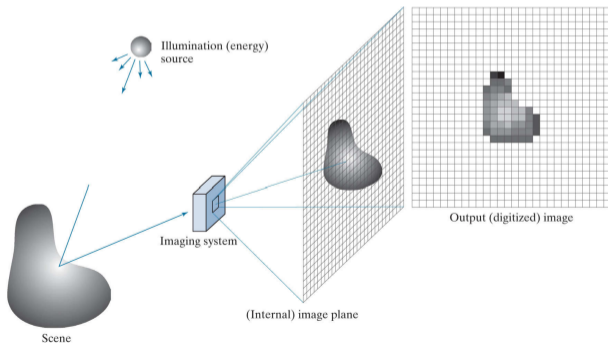
1.1. image acquisition

1. energy from an **illumination source** is reflected from a **scene**
2. the **imaging system** collects the incoming energy and focuses it onto an **image plane**
NB: light-sensing instruments typically use 2-D arrays of photosensors to record incoming light intensity $I(x)$: the CCD (*Charge-Coupled Device*)
3. the image plane is sampled and quantized to produce a **digital image**




1.1. image acquisition

1. energy from an **illumination source** is reflected from a **scene**
2. the **imaging system** collects the incoming energy and focuses it onto an **image plane**
NB: light-sensing instruments typically use 2-D arrays of photosensors to record incoming light intensity $I(x)$: the CCD (*Charge-Coupled Device*)
3. the image plane is sampled and quantized to produce a **digital image**

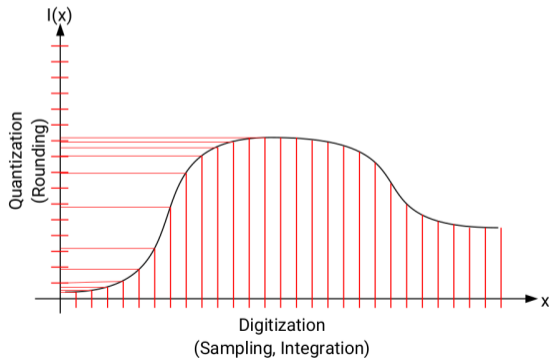


Credit: Gonzalez & Woods 2018

- each photosensor records incident light
- digitalization of an analog signal involves two operations
 - spatial sampling (= discretization of space domain)
 - intensity quantization (= discretization of incoming light signal)



- each photosensor records incident light
- digitalization of an analog signal involves two operations
 - **spatial sampling** (= discretization of space domain)
 - **intensity quantization** (= discretization of incoming light signal)



spatial sampling (= discretization of space domain)

⇒ smallest element resulting from the discretization of the space is called a pixel (=picture element)

(512, 512)



(128, 128)



(64, 64)



(32, 32)



intensity quantization (= discretization of light intensity signal)

⇒ typically, 256 levels (8 bits/pixel = 2^8 values) suffices to represent the intensity

8-bit resolution

$2^8 = 256$ gray levels



3-bit resolution

$2^3 = 8$ gray levels



2-bit resolution

$2^2 = 4$ gray levels



1-bit resolution

$2^1 = 2$ gray levels



spatial sampling (= discretization of space domain)

⇒ smallest element resulting from the discretization of the space is called a pixel (=picture element)

(512, 512)



(128, 128)



(64, 64)



(32, 32)



intensity quantization (= discretization of light intensity signal)

⇒ typically, 256 levels (8 bits/pixel = 2^8 values) suffices to represent the intensity

8-bit resolution

$2^8 = 256$ gray levels



3-bit resolution

$2^3 = 8$ gray levels



2-bit resolution

$2^2 = 4$ gray levels



1-bit resolution

$2^1 = 2$ gray levels



But how is the 3D world projected on a 2D plane?

⇒ comparison between human eye and pinhole camera:





Image = 3D world projection on 2D

⇒ projection using the **pinhole camera model**:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$


$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera** model:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$


$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera model**:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$


$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera model**:



Perspective transformation:

$$s m' = K[R|t]M' \quad (1)$$


$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

Image = 3D world projection on 2D

⇒ projection using the **pinhole camera model**:



Perspective transformation:

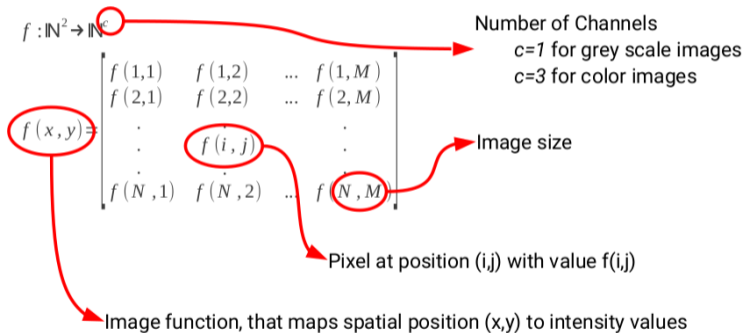
$$s m' = K[R|t]M' \quad (1)$$

$$s \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & u_0 \\ 0 & f_y & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_1 \\ r_{21} & r_{22} & r_{23} & t_2 \\ r_{31} & r_{32} & r_{33} & t_3 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix} \quad (2)$$

where:

- M' = 3D point in space with coordinates $[X, Y, Z]^T$ expressed in Euclidean coordinates
- m' = projection of the 3D point M' onto the image plane with coordinates $[u, v]^T$ expressed in pixel units
- K = camera calibration matrix (a.k.a. intrinsics parameters matrix)
 - f_x, f_y = focal lengths expressed in pixel units
 - u_0, v_0 = coordinates of the optical center (aka principal point), origin in the image plane
- $[R|t]$ = joint rotation-translation matrix (a.k.a. extrinsics parameters matrix), describing the camera pose, and translating from world coordinates to camera coordinates

⇒ digital image function $f(x, y)$



⇒ digital image function $f(x, y)$



	columns									
	0	1	2	3	4	5	6	7	8	9
0	24	24	67	103	87	79	176	138	94	180
1	98	53	66	226	34	34	241	240	24	143
2	228	107	60	58	144	251	137	93	86	130
3	155	108	132	159	129	141	245	211	100	24
4	91	187	67	135	49	175	193	61	24	183
5	199	251	80	24	121	105	222	147	226	63
6	181	27	56	238	113	158	176	47	167	109
7	24	172	24	192	184	162	181	202	24	72
8	24	106	30	24	53	68	178	232	91	219
9	211	181	78	24	24	185	204	106	131	70

Typical ranges:

- uint8 = [0-255]
(8 bits = 1 byte = $2^8 = 256$ values per pixel)
- float32 = [0-1]
(32 bits = 4 bytes = $4.3e9$ values per pixel)

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



(source [wikipedia](#))

How do we record colors?

⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene

2. Output of a 120×80 -pixel sensor with a Bayer filter


3. Output color-coded with Bayer filter colors

4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)

5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison

How do we record colors?


⇒ **Bayer Filter**: color filter array for arranging RGB color filters on a square grid of photosensors



1. Original scene
2. Output of a 120×80 -pixel sensor with a Bayer filter
3. Output color-coded with Bayer filter colors
4. Reconstructed image after interpolating missing color information (a.k.a. demosaicing)
5. Full RGB version at 120×80 -pixels for comparison


⇒ color image = 3D tensor in colorspace

- **RGB** = Red + Green + Blue bands (.JPEG)
- **RGBA** = Red + Green + Blue + Alpha bands (.PNG, .GIF, .BMP, TIFF, .JPEG 2000)




Other ways to represent the color information?

RGB colorspace


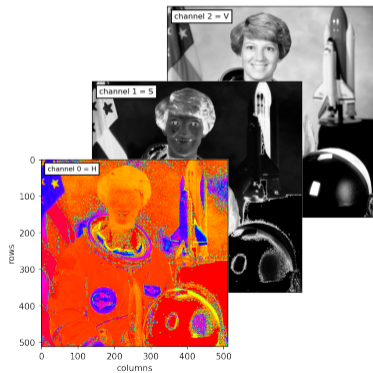


HSV colorspace



- Hue (H) = [0-360] \Rightarrow shift color
- Saturation (S) = [0-1] \Rightarrow shift intensity
- Value (V) = [0-1] \Rightarrow shift brightness

3D tensor with different information

RGB colorspaceHSV colorspace

- more saturation S

⇒ more intense colors



- more value V

⇒ brighter colors

- shift hue H

⇒ shift color

- more saturation S

⇒ more intense colors

original



saturation x2



original



value x1.5



- more value V

⇒ brighter colors

- shift hue H

⇒ shift color

- more saturation S

⇒ more intense colors

original



saturation x2



original



value x1.5



- more value V

⇒ brighter colors

original



hue x5




- shift hue H


⇒ shift color

Histogram of pixel values in each band:

original (uint8)




Histogram of pixel values after conversion from RGB (3-bands) to gray-scale (1-band):




Histogram of pixel values after conversion to float values (range [0-1])

gray-scale (float)




- original gray-scale




- histogram rescale to 10-90 percentiles
⇒ contrast stretching

- histogram equalize
⇒ spread out the most frequent intensity values

- original gray-scale




- histogram rescale to 10-90 percentiles
⇒ contrast stretching




- histogram equalize
⇒ spread out the most frequent intensity values


- original gray-scale



- histogram rescale to 10-90 percentiles
⇒ contrast stretching



- histogram equalize
⇒ spread out the most frequent intensity values



1. What is a digital image?


2. Point operations

1. homogeneous point operations
2. inhomogeneous Point Operations

3. Image processing levels

4. Image manipulation with Python


$$g(x, y) = T(f(x, y), x, y)$$



Inhomogeneous PO: T is dependent on (x,y)
Homogeneous PO: T is NOT dependent on (x,y)

Homogeneous Point Operations (does not depend on pixel position)

identity

 $g(x, y) = f(x, y)$ 

Homogeneous Point Operations (does not depend on pixel position)


identity



inverse



$$g(x, y) = 1 - f(x, y)$$



Homogeneous Point Operations (does not depend on pixel position)


identity



inverse



square root



Homogeneous Point Operations (does not depend on pixel position)

identity



inverse




square root



logarithm



$$g(x, y) = a \cdot \log(f(x, y) + 1)$$



Homogeneous Point Operations (does not depend on pixel position)

identity



inverse




square root



logarithm



exponential



Homogeneous Point Operations (does not depend on pixel position)

identity



inverse



square root




logarithm



exponential



square



Inhomogeneous Point Operations (depends on pixel position)


EX: background detection / change detection


 f_1

 f_i

 f_N


$$a(x, y) = \frac{1}{N} \sum_{i=0}^N f_i(x, y)$$

 $a(x, y)$


$$\begin{aligned} g_i(x, y) &= T(f(x, y), x, y) \\ &= f_i(x, y) - a(x, y) \end{aligned}$$

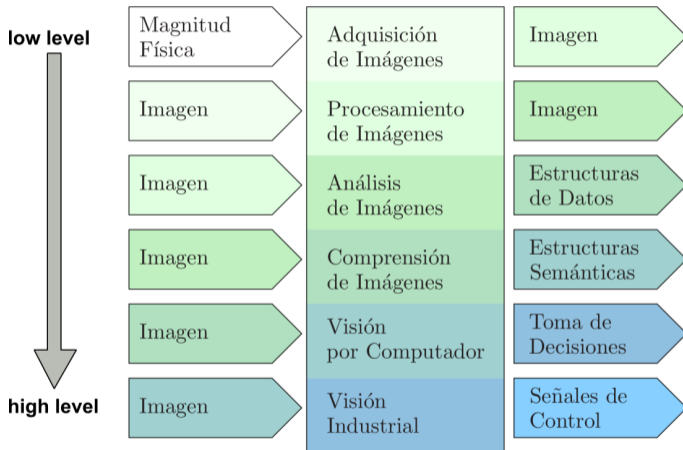
Inhomogeneous Point Operations (depends on pixel position)

EX: background detection / change detection



1. What is a digital image?
2. Point operations
3. Image processing levels
4. Image manipulation with Python

Image processing levels: inhomogeneous Point Operations



Credit: Pablo Alvarado 2012

Examples of processing levels:

- Low-level processing
 - image manipulation \Rightarrow *resizing, color adjustments, filtering, etc.*
 - feature extraction \Rightarrow *edges, gradients, etc.*
- Mid-level processing
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level processing
 - classification \Rightarrow *what is in the image?*
 - detection \Rightarrow *where are they?*
 - segmentation (semantic or instance) \Rightarrow *segment image and give names*

Examples of processing levels:

- Low-level processing
 - image manipulation \Rightarrow *resizing, color adjustments, filtering, etc.*
 - feature extraction \Rightarrow *edges, gradients, etc.*
- Mid-level processing
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level processing
 - classification \Rightarrow *what is in the image?*
 - detection \Rightarrow *where are they?*
 - segmentation (semantic or instance) \Rightarrow *segment image and give names*



hue x5




filter (high pass)




Examples of processing levels:


- Low-level processing
 - image manipulation \Rightarrow *resizing, color adjustments, filtering, etc.*
 - feature extraction \Rightarrow *edges, gradients, etc.*
- Mid-level processing
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level processing
 - classification \Rightarrow *what is in the image?*
 - detection \Rightarrow *where are they?*
 - segmentation (semantic or instance) \Rightarrow *segment image and give names*



Optical Flow (Farneback)

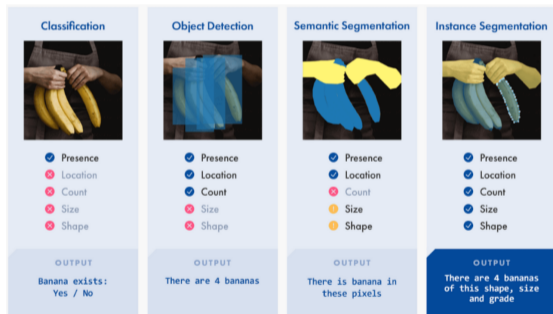


3D reconstruction



Examples of processing levels:

- Low-level processing
 - image manipulation \Rightarrow *resizing, color adjustments, filtering, etc.*
 - feature extraction \Rightarrow *edges, gradients, etc.*
- Mid-level processing
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level processing
 - classification \Rightarrow *what is in the image?*
 - detection \Rightarrow *where are they?*
 - segmentation (semantic or instance) \Rightarrow *segment image and give names*



Credit: cloudfactory

Examples of processing levels:

- Low-level processing
 - image manipulation \Rightarrow *resizing, color adjustments, filtering, etc.*
 - feature extraction \Rightarrow *edges, gradients, etc.*
- Mid-level processing
 - panorama stitching
 - Structure from Motion (SfM) \Rightarrow 2D to 3D
 - Optical Flow \Rightarrow velocities
- High-level processing
 - classification \Rightarrow *what is in the image?*
 - detection \Rightarrow *where are they?*
 - segmentation (semantic or instance) \Rightarrow *segment image and give names*

1. What is a digital image?
2. Point operations
3. Image processing levels
4. Image manipulation with Python
 1. numpy tutorial
 2. exercises

Numpy tutorial:

⇒ Open DIP4RS_02_imagebasics/[DIP4RS_02_numpy-tutorial.ipynb](#)

Exercices:

⇒ Open `DIP4RS_02_imagebasics/DIP4RS_02_exercices.ipynb`